

CUSTOMER MILESTONE BRIEF

MS #64 - CloudBank Credit Decision and AI Boundary

Purpose, customer value, delivered capability, evidence, and claim boundaries

Status	Complete - execution passed
Release	0.64.0
Date	2026-09-03
Audience	Customers, partners, executives, architects, delivery teams, and assurance reviewers

Published execution result

Published 2026-09-11: Eight-service target and edge controls. The original signed receipt is retained in the completed MS67 evidence chain.

[Original signed receipt 1](#)

[Full MS54–67 execution evidence index](#)

Executive summary

MS #64 - CloudBank Credit Decision and AI Boundary converts a specific part of the LIGHTYEAR modernization approach into a governed, reviewable capability. Completes the remaining CloudBank application workcells while making synthetic credit evidence and model-interaction controls explicit, testable, and resistant to overclaiming.

Why this matters to a customer

Completes the remaining CloudBank application workcells while making synthetic credit evidence and model-interaction controls explicit, testable, and resistant to overclaiming.

The practical result is a narrower, evidence-backed decision instead of a broad modernization claim. Commercial, architecture, delivery, security, and assurance teams can review the same capability, proof, and limits.

What the milestone delivers

- A signed MS #63 and MS #57 receipt chain requiring the same evidence key and PostgreSQL image.
- A fresh eight-service target with azn-server, checks, testrunner, creditscore, and chatbot explicitly marked migrated and each generated and executed in a dedicated Java workcell.
- An OAuth-protected Credit Score boundary with a distinct audience and a stable subject-and-UTC-date-bound HMAC result labelled synthetic-v1.
- A runtime-only minimum-length score pepper, no score persistence, and explicit rejection of real credit-bureau or regulated-model claims.

- A distinct-audience Chatbot boundary with system/user prompt separation, input and output limits, instruction-override rejection, and safe model failures.
- Per-authenticated-subject rate limiting and a host/TLS model-egress allowlist that permits plaintext only for allowlisted loopback endpoints.
- Eight executable Spring Boot JARs with zero Oracle and zero MicroTx runtime libraries.
- Twenty-eight bounded scenarios, twenty-one generated Java tests, safe aggregate diagnostics, deterministic contracts, schemas, cross-platform launchers, CI compilation and tests, Control Tower projection, documentation, and adversarial tests.

How it differs from earlier work

MS #57 qualified the PostgreSQL Customer workcell and MS #63 assembled Authorization, Account, Transfer, Checks, and Test Runner. MS #64 composes all prior targets, explicitly resolves the five remaining services as migrated, and generates and executes each of their target workcells in one eight-service target.

Earlier evidence is not automatically promoted into this milestone. A parser, simulator, local comparison, or generated artifact is not live platform equivalence or production readiness.

What a customer can do with it

A customer can use cloudbank credit decision and ai boundary as a bounded decision and delivery capability, attached to approved sources, owners, policies, target architecture, acceptance criteria, and authorized evidence.

Unresolved semantic loss, policy decisions, unsupported behavior, and live-evidence requirements remain visible.

Evidence and acceptance posture

The repository capability is complete because its committed implementation and release record are present. Customer qualification remains claim-specific and evidence-class-specific.

Review exact source and artifact identities, distinguish development from authorized live evidence, inspect policy and loss classifications, confirm passed and blocked gates, and verify that later changes have not invalidated the evidence.

Boundaries and claims not made

- These boundaries describe MS #64's own evidence scope. Later MS65, MS66 and MS67 receipts add their separate nonproduction qualifications; they do not rewrite this milestone's signed fields.
- Acceptance calls neither a credit bureau nor an external model and persists no prompts, responses, tokens, secrets, peppers, or upstream error details.
- A passing receipt qualifies generation and execution of all five remaining target workcells plus the bounded Credit Score and Chatbot application-control envelope and eight-service package, not a real credit decision or model-answer quality.
- Native Oracle AQ comparison, production data, whole-application Oracle/PostgreSQL equivalence, migration completion, production deployment, promotion, and production readiness remain false.

Source of record

- CHANGELOG.md release section(s): 0.64.0

- README.md product and operating guidance
- LIGHTYEAR-ROADMAP.md governing sequence and claim classifications
- docs/milestones/catalog.json metadata and docs/milestones/manifest.json artifact integrity